

Prithvi Shetty

Senior Data Engineer | BI Data Platforms | Snowflake | Databricks | Spark | AWS | Azure

prithvishetty096@gmail.com | +1 747-348-5817 | www.linkedin.com/in/prithvis319

PROFESSIONAL SUMMARY

Senior Data Engineer with 9+ years of experience building cloud-native data platforms, scalable ETL/ELT pipelines, BI-ready data models, and governed analytics systems across healthcare, insurance, banking, and financial services.

Deliver reliable data infrastructure using SQL, Python, Databricks, Spark, Snowflake, Azure, AWS, Airflow, Kafka, and DBT to improve data quality, data accessibility, business metrics, dashboard reliability, and downstream analytics readiness.

Bring strong experience translating business requirements into technical data solutions, modernizing regulated enterprise data environments, and building governed pipelines that support reporting, predictive analytics, and cross-functional decision-making.

SELECTED IMPACT

- Owned 20+ production data pipelines across healthcare, insurance, and financial services platforms.
- Integrated 25+ healthcare and insurance source systems into governed Azure, AWS, Databricks, Delta Lake, and Snowflake analytics platforms.
- Implemented 50+ validation, reconciliation, schema quality, and anomaly detection checks to improve downstream reporting trust.
- Processed millions of claims, transaction, clinical, operational, and third-party records across regulated enterprise data platforms.
- Reduced legacy ETL processing time by 30% through Informatica-to-Talend migration at Northern Trust.

TECHNICAL SKILLS

Cloud & Data Platforms: Azure, AWS, Snowflake, Databricks, Azure Data Factory, ADLS Gen2, Azure Synapse, Azure SQL Database, Azure Stream Analytics, Azure DevOps, Azure Monitor, AWS Glue, AWS Lambda, S3, Redshift, EMR, RDS, IAM, KMS, CloudWatch, AWS Data Migration Service

Data Engineering & Orchestration: SQL, Python, PySpark, Apache Spark, Spark Structured Streaming, Delta Lake, ETL/ELT, Apache Airflow, Apache Kafka, Apache NiFi, Informatica, Talend, Hadoop, Hive, REST API ingestion, Parquet, JSON, XML, fixed-width files, delimited files, structured data, semi-structured data

BI, Modeling & Analytics: Snowflake, AWS Redshift, Azure SQL Database, PostgreSQL, Star Schema, Fact and Dimension Modeling, Logical and Physical Data Modeling, Data Warehouse Design, Integration Tables, Incremental Pipelines, DBT, Source-to-Target Mapping, Data Contracts, Tableau, Power BI

Governance, Quality & Observability: Unity Catalog, Data Lineage, Schema Validation, Reconciliation Checks, Data Quality Controls, Anomaly Detection, HIPAA Standards, Encryption, Azure Key Vault, Azure Active Directory, AWS Security Hub, Azure Monitor, Log Analytics, AWS CloudWatch, Custom Metrics, Alerting, SLA Monitoring

DevOps & Delivery: Terraform, Jenkins, Git, Azure DevOps, Docker, Kubernetes, AKS, EKS, CI/CD, Deployment Automation, Monitoring Pipelines, Release Consistency, Technical Documentation, Agile/Scrum Collaboration

AI/ML Data Enablement: MLflow, Feature Engineering, Predictive Analytics, ML-Ready Data Assets, AI/ML Data Pipelines

PROFESSIONAL EXPERIENCE

Client: Galderma, New York, NY | Nov 2023 - Present

Role: Senior Data Engineer

Project: Healthcare Data Platform Modernization & Predictive Analytics

- Owned 20+ production healthcare data pipelines across Azure, Databricks, and Snowflake to support governed enterprise analytics and business decision-making.
- Integrated 10+ clinical, claims, operational, and third-party source systems into lakehouse and warehouse data models supporting reporting, analytics, and downstream data accessibility.
- Built Delta Lake and PySpark processing patterns for terabyte-scale healthcare data used in BI reporting, predictive analytics, and ML-ready workflows.
- Designed Snowflake and DBT analytical models using fact and dimension structures to support Tableau, Power BI, and decision-support datasets.
- Created curated integration-ready datasets across clinical, claims, operational, and third-party data sources to improve consistency of business logic across reporting layers.
- Implemented Unity Catalog governance, lineage, access controls, and Key Vault-backed security to support HIPAA-aligned data protection and enterprise data standards.
- Added 50+ validation, reconciliation, schema quality, and anomaly detection checks to improve data quality and reduce downstream reporting risk.
- Improved Databricks and Spark job performance by 25-45% through partitioning, caching, DAG tuning, and workload optimization across high-volume workflows.
- Developed real-time ingestion and transformation workflows using Azure Stream Analytics and Spark Structured Streaming for continuous healthcare data monitoring.
- Automated data workflows using Apache Airflow, Azure Data Factory, and CI/CD practices to improve processing reliability, reproducibility, and release consistency.
- Developed REST API ingestion frameworks using Python and Airflow to onboard third-party healthcare data sources into governed ETL workflows.
- Partnered with analytics, data science, architecture, and business stakeholders to translate reporting and predictive analytics requirements into scalable data models and pipelines.
- Built feature engineering and ML data workflows in Databricks to support predictive analytics and AI-driven healthcare insights.
- Delivered governed Tableau and Power BI reporting layers supporting 50+ clinical, operational, and leadership stakeholders.

Client: BCBS, Chicago, IL | Dec 2021 - Nov 2023

Role: Data Engineer

Project: Insurance Data Platform Modernization & Analytics

- Built and supported 30+ production insurance data pipelines across batch, streaming, and orchestration workflows for enterprise reporting and analytics.
- Integrated 15+ claims, policy, transactional, and operational source systems into AWS, Snowflake, and Redshift analytics layers.
- Developed Spark processing workloads on AWS EMR using Python and SQL to process millions of claims and transaction records for insurance analytics.
- Designed AWS Redshift warehouse structures using distribution keys, sort keys, and columnar storage to support scalable business intelligence and reporting.
- Improved Redshift reporting query performance by 35% through workload tuning, table design optimization, and data warehouse performance improvements.
- Defined source-to-target mappings, data contracts, reusable transformation layers, and standardized business logic across Snowflake and Redshift environments.
- Built analytics-ready integration datasets for risk scoring, claims forecasting, underwriting optimization, operational reporting, and downstream BI use cases.
- Maintained 98%+ SLA adherence across scheduled insurance data pipelines through monitoring, alerting, failure detection, and operational support workflows.

- Implemented Apache Airflow orchestration to automate batch ETL workflows, manage dependencies, and improve timely data availability for reporting.
- Established AWS security controls through access policies, encryption patterns, and monitoring for regulated insurance datasets.
- Led cloud migration initiatives that consolidated legacy insurance systems while preserving data integrity and reporting continuity.

Client: Northern Trust Bank, Chicago, IL | May 2020 - Dec 2021

Role: ETL Engineer

Project: Financial Data Warehouse Modernization & Real-Time Analytics

- Maintained and modernized 25+ Informatica and Talend ETL workflows supporting transaction, trade, and risk reporting datasets.
- Consolidated transaction, trade, and risk data into 15+ curated reporting datasets used for financial analytics and regulatory reporting.
- Built cloud ETL workflows using AWS Glue, Lambda, S3, and Redshift for daily financial reporting, analytics delivery, and operational data processing.
- Developed Kafka-based streaming workflows to support near real-time trading activity, regulatory monitoring, and operational reporting.
- Designed AWS Redshift warehouse models using partitioning, distribution keys, and sort keys to support complex financial queries and business reporting.
- Implemented 20-50 validation and reconciliation checks to improve data accuracy across trade, transaction, and risk reporting datasets.
- Automated ETL orchestration using Apache Airflow to improve dependency management, batch reliability, and downstream data availability.
- Optimized Hadoop and Spark workloads on AWS EMR through DAG tuning, partitioning, and caching for large-scale batch processing.
- Led migration of legacy ETL workflows to Talend, improving data reliability and reducing processing time by 30%.
- Built CI/CD release patterns and CloudWatch monitoring workflows to improve deployment consistency, pipeline visibility, and operational reliability.

Client: Luxoft India Pvt. Ltd., India | Jul 2017 - Jul 2018

Role: Hadoop Engineer

- Administered Hadoop clusters supporting scalable storage, batch processing, and distributed data workloads for enterprise datasets.
- Developed Hive, Spark, SQL, and Python ETL pipelines to ingest, transform, and integrate structured and semi-structured datasets.
- Tuned 20+ Hadoop and Spark jobs using partitioning, compression, and query optimization to improve throughput and reduce resource contention.
- Automated build and deployment workflows through Jenkins and supported Docker-based workloads in Linux environments.

Client: Optimum Info System Pvt. Ltd., India | Nov 2015 - Jun 2017

Role: Data Analyst

- Automated data extraction, cleaning, and transformation workflows using Python, pandas, NumPy, and SQL for analytics-ready datasets.
- Built 10+ Tableau dashboards to visualize KPIs, improve business visibility, and support operational decision-making.
- Optimized SQL queries, implemented validation checks, and translated business requirements into structured reporting workflows.

EDUCATION

California State University Northridge

Master of Science in Engineering Management | 2018 - 2020

CERTIFICATIONS

Microsoft Certified: Azure Data Engineer Associate

Databricks Certified Data Engineer Associate